

# SVM Regression on Boston Housing Data

*Karen Mazidi*

## Load the data

Read more about the data by typing “?Boston” at the console.

```
library(MASS)
df <- Boston[]
str(df)

## 'data.frame':  506 obs. of  14 variables:
## $ crim   : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn     : num  18 0 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus  : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas   : int   0 0 0 0 0 0 0 0 0 0 ...
## $ nox    : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm     : num  6.58 6.42 7.18 7 7.15 ...
## $ age    : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis    : num  4.09 4.97 4.97 6.06 6.06 ...
## $ rad    : int   1 2 2 3 3 3 5 5 5 5 ...
## $ tax    : num  296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num  15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ black  : num  397 397 393 395 397 ...
## $ lstat  : num  4.98 9.14 4.03 2.94 5.33 ...
## $ medv   : num  24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

## Train and test

Divide the data into 80% train and 20% test.

```
set.seed(1234)
i <- sample(nrow(Boston), 0.8*nrow(Boston), replace=FALSE)
train <- df[i,]
test <- df[-i,]
```

## Linear regression

Build a linear regression model on the training data.

```
lm1 <- lm(medv~., data=train)
summary(lm1)

##
## Call:
## lm(formula = medv ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -13.537  -2.913  -0.546   1.848  24.915
##
## Coefficients:
```

```
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) 44.900577   6.016980   7.462 5.59e-13 ***
## crim       -0.085000   0.049892  -1.704 0.08924 .
## zn         0.047219   0.015849   2.979 0.00307 **
## indus      0.038249   0.070942   0.539 0.59008
## chas       2.724575   0.966685   2.818 0.00507 **
## nox       -19.139048   4.382515  -4.367 1.62e-05 ***
## rm         2.949428   0.479982   6.145 1.98e-09 ***
## age       -0.007757   0.015670  -0.495 0.62087
## dis       -1.558391   0.224867  -6.930 1.75e-11 ***
## rad        0.302988   0.076673   3.952 9.21e-05 ***
## tax       -0.012284   0.004206  -2.920 0.00370 **
## ptratio   -1.008491   0.152951  -6.594 1.40e-10 ***
## black      0.008717   0.003345   2.606 0.00951 **
## lstat     -0.555420   0.056482  -9.834 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.884 on 390 degrees of freedom
## Multiple R-squared:  0.7203, Adjusted R-squared:  0.711
## F-statistic: 77.28 on 13 and 390 DF, p-value: < 2.2e-16
```

## Evaluate on the test data

We have 90% correlation between the predicted and target home prices. The mse of 18.9 is a rmse of 4.347, so we are off by about \$4,347 on average for the homes in the neighborhood. That's pretty good.

```
pred_lm <- predict(lm1, newdata=test)
cor(pred_lm, test$medv)
```

```
## [1] 0.900081
```

```
mse_lm <- mean((pred_lm - test$medv)^2)
mse_lm
```

```
## [1] 18.93611
```

## SVM Linear

Now we try SVM regression with a linear kernel.

```
library(e1071)
svm_fit1 <- svm(medv~., data=train, kernel="linear", cost=10, scale=FALSE)
summary(svm_fit1)
```

```
##
## Call:
## svm(formula = medv ~ ., data = train, kernel = "linear", cost = 10,
##      scale = FALSE)
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: linear
##      cost:   10
```

```
##      gamma: 0.07692308
##      epsilon: 0.1
##
##
## Number of Support Vectors: 392
svm_pred1 <- predict(svm_fit1, newdata=test)
cor(svm_pred1, test$medv)

## [1] 0.912717
mse_svm1 <- mean((svm_pred1 - test$medv)^2)
mse_svm1

## [1] 18.4258
```

## The tune() function

The linear SVM did slightly better than linear regression. Next we perform some tuning to find the best cost parameter. The tune() function tries to find optimal hyperparameters for the svm using a grid search. This involves trying all the suggested parameters in a cross-validation scheme. Here we asked it to try several different cost parameters. The best model can be extracted from the tune results.

```
tune_svm1 <- tune(svm, medv~., data=train, kernel="linear",
                 ranges=list(cost=c(0.001, 0.01, 0.1, 1, 5, 10, 100)))
summary(tune_svm1)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.1
##
## - best performance: 26.97133
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 44.80447    17.51189
## 2 1e-02 28.50440    12.25274
## 3 1e-01 26.97133    11.04768
## 4 1e+00 27.07120    11.18866
## 5 5e+00 27.07697    11.17431
## 6 1e+01 27.08590    11.19209
## 7 1e+02 27.06015    11.16789

best_mod1 <- tune_svm1$best.model
summary(best_mod1)

##
## Call:
## best.tune(method = svm, train.x = medv ~ ., data = train, ranges = list(cost = c(0.001,
##   0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
##
```

```
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: linear
##       cost:  0.1
##       gamma: 0.07692308
##   epsilon:  0.1
##
##
## Number of Support Vectors:  318
```

## Use the best model

The best model parameters were selected on the train set. It will not necessarily perform better on the test data, and indeed it performed slightly worse.

```
svm_pred2 <- predict(best_mod1, newdata=test)
cor(svm_pred2, test$medv)
```

```
## [1] 0.9074539
```

```
mse_svm2 <- mean((svm_pred2 - test$medv)^2)
mse_svm2
```

```
## [1] 19.36801
```

## Try the radial kernel

For radial kernel we have an additional hyperparameter, gamma.

```
svm_fit2 <- svm(medv~., data=train, kernel="radial", cost=1, gamma=1, scale=FALSE)
summary(svm_fit2)
```

```
##
## Call:
## svm(formula = medv ~ ., data = train, kernel = "radial", cost = 1,
##     gamma = 1, scale = FALSE)
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: radial
##       cost:  1
##       gamma:  1
##   epsilon:  0.1
##
##
## Number of Support Vectors:  398
```

```
svm_pred2 <- predict(svm_fit2, newdata=test)
cor(svm_pred2, test$medv)
```

```
## [1] -0.01849452
```

```
mse_svm2 <- mean((svm_pred2 - test$medv)^2)
mse_svm2
```

```
## [1] 95.05349
```

### Tune the hyperparameters

```
set.seed(1234)
tune_svm2 = tune(svm, medv~., data=train, kernel="radial",
                 ranges=list(cost=c(0.1,1,10,100,1000),
                             gamma=c(0.5,1,2,3,4)))
summary(tune_svm2)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##    10   0.5
##
## - best performance: 22.04514
##
## - Detailed performance results:
##   cost gamma   error dispersion
## 1  1e-01   0.5 58.00975 26.042036
## 2  1e+00   0.5 28.57541 13.694933
## 3  1e+01   0.5 22.04514  8.732259
## 4  1e+02   0.5 22.09964  8.574124
## 5  1e+03   0.5 22.09964  8.574124
## 6  1e-01   1.0 68.63374 28.598236
## 7  1e+00   1.0 39.62097 20.094040
## 8  1e+01   1.0 33.87235 15.209260
## 9  1e+02   1.0 33.87235 15.209260
## 10 1e+03   1.0 33.87235 15.209260
## 11 1e-01   2.0 76.36623 29.667097
## 12 1e+00   2.0 53.73030 26.371605
## 13 1e+01   2.0 48.55734 22.109707
## 14 1e+02   2.0 48.55734 22.109707
## 15 1e+03   2.0 48.55734 22.109707
## 16 1e-01   3.0 79.27019 29.696176
## 17 1e+00   3.0 61.93589 28.314759
## 18 1e+01   3.0 57.29582 25.017593
## 19 1e+02   3.0 57.29582 25.017593
## 20 1e+03   3.0 57.29582 25.017593
## 21 1e-01   4.0 80.74464 29.632789
## 22 1e+00   4.0 66.95143 28.820064
## 23 1e+01   4.0 62.79526 26.044450
## 24 1e+02   4.0 62.79526 26.044450
## 25 1e+03   4.0 62.79526 26.044450
```

Again, best model isn't.

```
svm_pred3 <- predict(tune_svm2$best.model, newdata=test)
cor(svm_pred3, test$medv)
```

```
## [1] 0.8926458
```

```
mse_svm3 <- mean((svm_pred3 - test$medv)^2)
mse_svm3
```

```
## [1] 19.41668
```