

Text Processing with tm

Karen Mazidi

Amazon review data from Kaggle: <https://www.kaggle.com/bittlingmayer/amazonreviews>

Load data and package tm

```
library(tm)
```

```
## Loading required package: NLP
```

```
reviews <- read.csv("data/reviews.csv", header=TRUE, stringsAsFactors=F)
```

counts for ratings

Rating = 1 means 1 or 2 stars; Rate=2 mean 4 or 5 stars; 3 stars were ignored

About 46% were low and 54% high.

```
low_ratings <- nrow(reviews[reviews$Rating == 1,])  
high_ratings <- nrow(reviews[reviews$Rating == 2,])
```

Make a simple corpus

```
am_corpus <- Corpus(VectorSource(reviews$Review))  
inspect(am_corpus[1])
```

```
## <<SimpleCorpus>>
```

```
## Metadata: corpus specific: 1, document level (indexed): 0
```

```
## Content: documents: 1
```

```
##
```

```
## [1] Stuning even for the non-gamer: This sound track was beautiful! It paints the senery in your mind
```

Preprocess

```
am_clean <- tm_map(am_corpus, content_transformer(tolower))  
am_clean <- tm_map(am_clean, removeNumbers)  
am_clean <- tm_map(am_clean, removePunctuation)  
am_clean <- tm_map(am_clean, removeWords, stopwords())  
am_clean <- tm_map(am_clean, stripWhitespace)
```

Make Document Term Matrix

```
am_dtm <- DocumentTermMatrix(am_clean)  
am_dtm
```

```
## <<DocumentTermMatrix (documents: 4139, terms: 21669)>>
## Non-/sparse entries: 143646/89544345
## Sparsity          : 100%
## Maximal term length: 106
## Weighting          : term frequency (tf)
```

Divide into test and train

```
set.seed(1234)
i <- sample(nrow(reviews), 0.75*nrow(reviews), replace=FALSE)
# labels
train_labels <- reviews[i, 1]
test_labels <- reviews[-i, 1]
# data
train <- am_clean[i]
test <- am_clean[-i]
```

Ignore rare words

```
freq_words <- findFreqTerms(am_dtm, 5)
train <- DocumentTermMatrix(train, control=list(dictionary=freq_words))
test <- DocumentTermMatrix(test, control=list(dictionary=freq_words))
inspect(train[50:55, 200:209])
```

```
## <<DocumentTermMatrix (documents: 6, terms: 10)>>
## Non-/sparse entries: 3/57
## Sparsity          : 95%
## Maximal term length: 7
## Weighting          : term frequency (tf)
## Sample            :
##      Terms
## Docs bought send someone spent term thing unless want worst written
## 50      0    0      0    0    0    0    0    0    0    0
## 51      0    1      0    0    0    0    0    0    0    0
## 52      0    2      0    0    0    0    0    0    0    0
## 53      0    1      0    0    0    0    0    0    0    0
## 54      0    0      0    0    0    0    0    0    0    0
## 55      0    0      0    0    0    0    0    0    0    0
```

Create binary matrix

```
convert_count <- function(x) {
  y <- ifelse(x>0, 1, 0)
  y <- factor(y)
  y
}

train <- apply(train, 2, convert_count)
test <- apply(test, 2, convert_count)
```

Naive Bayes

```
library(e1071)
nb1 <- naiveBayes(train, factor(train_labels))
```

Evaluate on test data

We only got 54% accuracy. Further improvements might be made by removing proper nouns and stemming, but we leave that for future work.

```
pred <- predict(nb1, newdata=test)
table(pred, test_labels)
```

```
##      test_labels
## pred    1    2
##      1 303 233
##      2 239 260
```

```
mean(pred == test_labels)
```

```
## [1] 0.5439614
```